

Concept Learning

TM Chapter 2

Outline

- Learning from examples
- General-to-specific ordering over hypotheses
- Version spaces and candidate elimination algorithm
- Picking new examples
- The need for inductive bias

Why Study Concept Learning?

- Simple (e.g., **assumes error-free, noise-free data**) and practically useless!
- Easier to fully understand and explain the **general challenges, issues, and concepts in ML**
- **White-box** model: Prediction is interpretable & explainable
- Relates well to your earlier modules: Discrete math (CS1231), logic, search (CS3243)
- Principled and rigorously grounded: Proofs, proofs, and more proofs!

What is Concept Learning?

- **What is a concept?** A boolean-valued function over a set of input instances (each comprising input attributes)
- **Concept learning is a form of supervised learning.** Infer an unknown boolean-valued function from **training examples**

Example	<i>Sky</i>	<i>AirTemp</i>	<i>Humidity</i>	<i>Wind</i>	<i>Water</i>	<i>Forecast</i>	<i>EnjoySport</i>
1	Sunny	Warm	Normal	Strong	Warm	Same	Yes
2	Sunny	Warm	High	Strong	Warm	Same	Yes
3	Rainy	Cold	High	Strong	Warm	Change	No
4	Sunny	Warm	High	Strong	Cool	Change	Yes

Positive (+ve) and negative (–ve) **training examples** for target concept *EnjoySport*

How to Represent a Hypothesis?

- Many possible hypothesis representations: Trade-off between **expressive power** vs. **smaller hypothesis space**
- Consider a simple representation: Hypothesis h is a **conjunction of constraints on input attributes**
- Each **constraint** can be
 - a specific value (e.g., “*Water = Warm*”)
 - don’t care (e.g., “*Water = ?*”)
 - no value allowed (e.g., “*Water = \emptyset* ”)

<i>Sky</i>	<i>AirTemp</i>	<i>Humidity</i>	<i>Wind</i>	<i>Water</i>	<i>Forecast</i>
<i>⟨Sunny,</i>	<i>?,</i>	<i>?,</i>	<i>Strong,</i>	<i>?,</i>	<i>Same⟩</i>

Concept Learning for *EnjoySport*

Given

- **Input instances** X : Each instance $x \in X$ is represented by the following input attributes describing the day:
 - *Sky* (with possible values *Sunny*, *Cloudy*, and *Rainy*)
 - *AirTemp* (with values *Warm* and *Cold*)
 - *Humidity* (with values *Normal* and *High*)
 - *Wind* (with values *Strong* and *Weak*)
 - *Water* (with values *Warm* and *Cool*)
 - *Forecast* (with values *Same* and *Change*)

Concept Learning for *EnjoySport*

Given

- **Hypothesis space** H : Each hypothesis $h \in H$ ($h : X \rightarrow \{0, 1\}$) is represented by a conjunction of constraints (see page 5) on input attributes (e.g., $\langle \text{Sunny}, ?, ?, \text{Strong}, ?, \text{Same} \rangle$).

Definition. An input instance $x \in X$ **satisfies** (all constraints of) a hypothesis $h \in H$ iff $h(x) = 1$.

In other words, h classifies x as a +ve example.

Concept Learning for *EnjoySport*

Given

- Unknown **target concept/function** *EnjoySport*: $c : X \rightarrow \{0, 1\}$
- Noise-free **training examples** D : +ve and -ve examples of the target function (e.g., $D = \{\langle x_k, c(x_k) \rangle\}_{k=1, \dots, n}$)

Determine a hypothesis $h \in H$ that is **consistent** with D

Definition. A hypothesis h is **consistent** with a set of training examples D iff $h(x) = c(x)$ for all $\langle x, c(x) \rangle \in D$.

How is saying ‘ h is **consistent** with training example $\langle x, c(x) \rangle$ ’ different from ‘saying instance x **satisfies** h ’? Implication?

Inductive Learning Assumption

Any hypothesis found to approximate the target function well over a sufficiently large set of **training** examples will also approximate the target function well over other **unobserved** examples.

Concept Learning is Search

Goal. Search for a hypothesis $h \in H$ that is **consistent** with D

For *EnjoySport* task, H contains

- $5 \times 4 \times 4 \times 4 \times 4 \times 4 = 5120$ syntactically distinct hypotheses
- $1 + 4 \times 3 \times 3 \times 3 \times 3 \times 3 = 973$ semantically distinct hypotheses
- Every hypothesis containing 1 or more \emptyset symbols represents an empty set of input instances, hence classifying every instance as a –ve example


In practice, hypothesis space H is much **larger** and possibly infinite, hence motivating the need to **exploit structure** for searching efficiently.

Exploit Structure in Concept Learning

Definition. h_j is **more general than or equal to** h_k (denoted by $h_j \geq_g h_k$) iff any input instance x that satisfies h_k also satisfies h_j :

$$\forall x \in X \ (h_k(x) = 1) \rightarrow (h_j(x) = 1) .$$

\geq_g relation defines a **partial order** (reflexive, antisymmetric, transitive) over H & not total order (e.g., $h_1 \not\geq_g h_3$ and $h_3 \not\geq_g h_1$)



Exploit Structure in Concept Learning

Definition. h_j is **more general than or equal to** h_k (denoted by $h_j \geq_g h_k$) iff any input instance x that satisfies h_k also satisfies h_j :

$$\forall x \in X \ (h_k(x) = 1) \rightarrow (h_j(x) = 1) .$$

Definition. h_j is **more general than** h_k (denoted by $h_j >_g h_k$) iff $h_j \geq_g h_k$ and $h_k \not\geq_g h_j$.

Definition. h_j is **more specific than** h_k iff h_k is **more general than** h_j .

Definition. h_j is **more specific than or equal to** h_k iff h_k is **more general than or equal to** h_j .


FIND-S Algorithm

Idea. Start with most specific hypothesis. Whenever it wrongly classifies a +ve training example as -ve, “minimally” generalize it to satisfy its input instance.

1. Initialize h to most specific hypothesis in H
2. For each positive training instance x
 - For each attribute constraint a_i in h
 - If x satisfies constraint a_i in h
 - Then do nothing
 - Else replace a_i in h by the next more general constraint that is satisfied by x
3. Output hypothesis h


Hypothesis Space Search by FIND-S

Proposition 1. h is consistent with D iff every +ve training instance satisfies h and every -ve training instance does not satisfy h .



Hypothesis Space Search by FIND-S

Proposition 2. Suppose that $c \in H$. Then, h_n is consistent with $D = \{\langle x_k, c(x_k) \rangle\}_{k=1, \dots, n}$.



Limitations of FIND-S

- Can't tell whether Find-S has learned target concept
- Can't tell when training examples are inconsistent (i.e., contain **errors or noise**)
- Picks a maximally specific h (why?)
- Depending on H , there might be several!

Version Spaces

Definition. The **version space** $VS_{H,D}$ wrt hypothesis space H and training examples D , is the subset of hypotheses from H consistent with D :

$$VS_{H,D} = \{h \in H \mid h \text{ is consistent with } D\} .$$

- If $c \in H$, then a large enough D can reduce $VS_{H,D}$ to $\{c\}$
- If D is insufficient, then $VS_{H,D}$ represents the **uncertainty** of what the target concept is
- $VS_{H,D}$ contains all consistent hypotheses, including maximally specific hypotheses

LIST-THEN-ELIMINATE Algorithm

Idea. List all hypotheses in H . Then, eliminate any hypothesis found inconsistent with any training example.

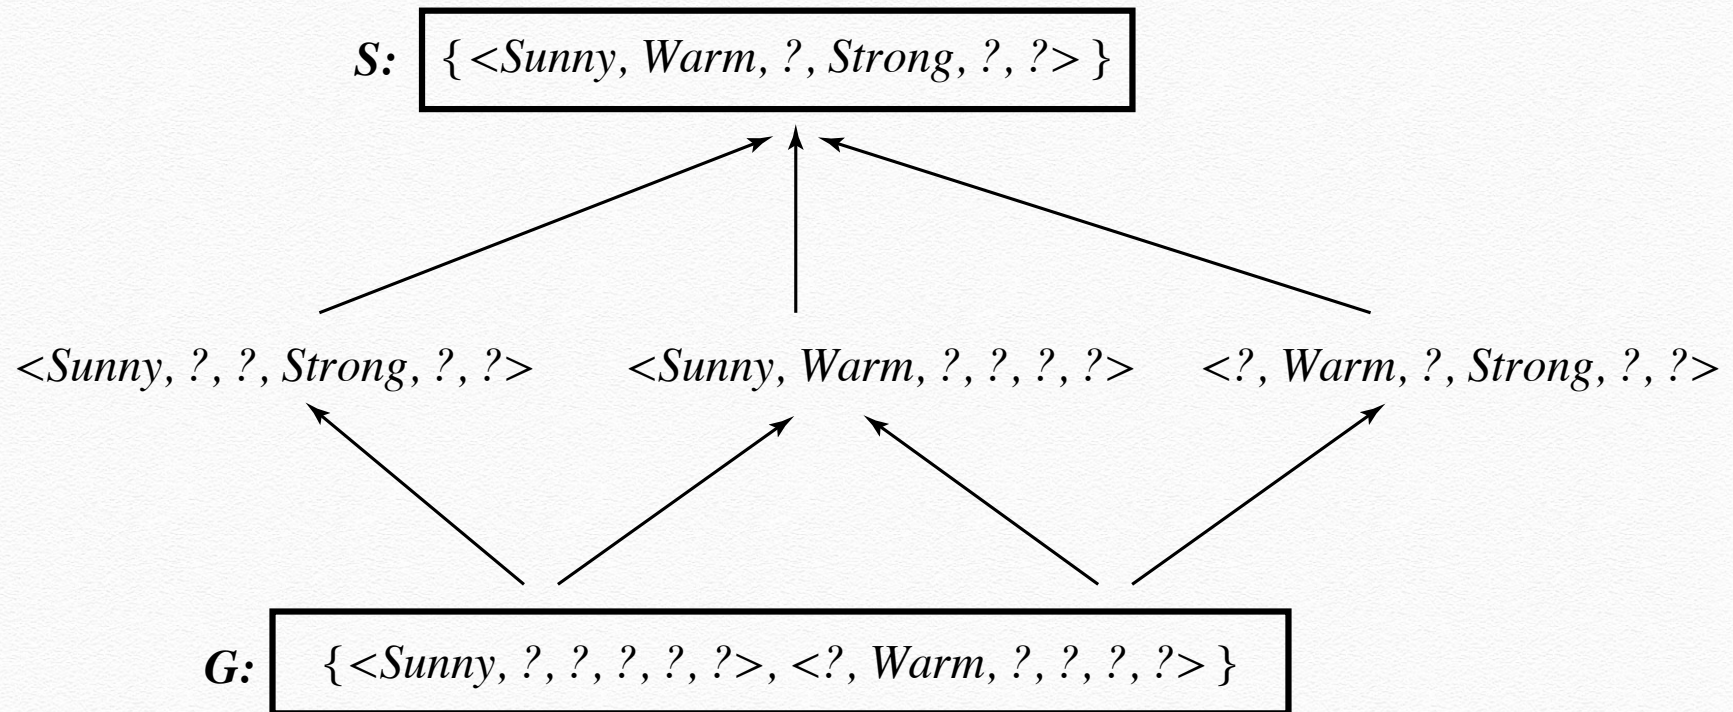
1. $VersionSpace \leftarrow$ a list containing every hypothesis in H
2. For each training example $\langle x, c(x) \rangle$

Remove from $VersionSpace$ any hypothesis h for which
 $h(x) \neq c(x)$

3. Output the list of hypotheses in $VersionSpace$

Limitation. Prohibitively expensive to exhaustively enumerate all hypotheses in finite H

Version Space for *EnjoySport*



Compact Representation of Version Space

Definition. The **general boundary** G of $VS_{H,D}$ is the set of maximally general members of H consistent with D :

$$G = \{g \in H \mid g \text{ consistent with } D \wedge (\neg \exists g' \in H \ g' >_g g \wedge g' \text{ consistent with } D)\}.$$

Definition. The **specific boundary** S of $VS_{H,D}$ is the set of maximally specific members of H consistent with D :

$$S = \{s \in H \mid s \text{ consistent with } D \wedge (\neg \exists s' \in H \ s >_g s' \wedge s' \text{ consistent with } D)\}.$$

Every member of version space lies between these boundaries:

Version space representation theorem (VSRT).

$$VS_{H,D} = \{h \in H \mid \exists s \in S \ \exists g \in G \ g \geq_g h \geq_g s\} .$$

Proof of Version Space Representation Theorem

\Leftarrow Every h satisfying RHS is in $VS_{H,D}$.

1. Choose arbitrary $g \in G, s \in S, h \in H$ s.t. $g \succeq_g h \succeq_g s$
2. Every +ve training instance satisfies s , by Def. of S and Prop. 1
3. Since $h \succeq_g s$, every +ve training instance satisfies h
4. Every -ve training instance does not satisfy g , by Def. of G and Prop. 1
5. Since $g \succeq_g h$, every -ve training instance does not satisfy h
6. h is consistent with D , by Prop. 1 and steps 3 and 5
7. $h \in VS_{H,D}$

\Rightarrow Every member of $VS_{H,D}$ satisfies RHS. DIY.

CANDIDATE-ELIMINATION Algorithm

Idea. Start with most general and specific hypotheses. Each training example “minimally” generalizes S and specializes G to remove inconsistent hypotheses from version space.

1. $G \leftarrow$ maximally general hypotheses in H
2. $S \leftarrow$ maximally specific hypotheses in H

CANDIDATE-ELIMINATION Algorithm

3. For each training example d

- If d is a +ve example
 - Remove from G any hypothesis inconsistent with d
 - For each $s \in S$ not consistent with d
 - ▶ Remove s from S
 - ▶ Add to S all minimal generalizations h of s s.t.
 h is consistent with d , and
some member of G is more general than or
equal to h
 - ▶ Remove from S any hypothesis that is more
general than another hypothesis in S

CANDIDATE-ELIMINATION Algorithm

/* Influence of +ve and -ve examples on S and G are dual */

- If d is a -ve example
 - Remove from S any hypothesis inconsistent with d
 - For each $g \in G$ not consistent with d
 - ▶ Remove g from G
 - ▶ Add to G all minimal specializations h of g s.t.
 h is consistent with d , and
some member of S is more specific than or equal to h
 - ▶ Remove from G any hypothesis that is more specific than another hypothesis in G

CANDIDATE-ELIMINATION Trace 1

- Remove from G any hypothesis inconsistent with d
- For each $s \in S$ not consistent with d

- ▶ Remove s from S
- ▶ Add to S all minimal generalizations h of s s.t.
 h is consistent with d , and
 some member of G is more general than or
 equal to h

S_0 : $\{\langle \emptyset, \emptyset, \emptyset, \emptyset, \emptyset, \emptyset \rangle\}$

S_1 : $\{\langle \text{Sunny}, \text{Warm}, \text{Normal}, \text{Strong}, \text{Warm}, \text{Same} \rangle\}$

S_2 : $\{\langle \text{Sunny}, \text{Warm}, ?, \text{Strong}, \text{Warm}, \text{Same} \rangle\}$

G_0, G_1, G_2 : $\{\langle ?, ?, ?, ?, ?, ? \rangle\}$

Training examples:

1. $\langle \text{Sunny}, \text{Warm}, \text{Normal}, \text{Strong}, \text{Warm}, \text{Same} \rangle, \text{Enjoy Sport} = \text{Yes}$
2. $\langle \text{Sunny}, \text{Warm}, \text{High}, \text{Strong}, \text{Warm}, \text{Same} \rangle, \text{Enjoy Sport} = \text{Yes}$


CANDIDATE-ELIMINATION Trace 2

- Remove from S any hypothesis inconsistent with d
- For each $g \in G$ not consistent with d
 - ▶ Remove g from G
 - ▶ Add to G all minimal specializations h of g s.t.

h is consistent with d , and

$S_2, S_3: \{ \langle \text{Sunny}, \text{Warm}, ?, \text{Strong}, \text{Warm}, \text{Same} \rangle \}$

some member of S is more specific than or equal to h



Training Example:

3. $\langle \text{Rainy}, \text{Cold}, \text{High}, \text{Strong}, \text{Warm}, \text{Change} \rangle, \text{EnjoySport} = \text{No}$

CANDIDATE-ELIMINATION Trace 3

- Remove from G any hypothesis inconsistent with d
- For each $s \in S$ not consistent with d
 - ▶ Remove s from S
 - ▶ Add to S all minimal generalizations h of s s.t.

h is consistent with d , and

some member of G is more general than or equal to h

S_3 : {<Sunny, Warm, ?, Strong, Warm, Same>}



S_4 : {<Sunny, Warm, ?, Strong, ?, ?>}

G_4 : {<Sunny, ?, ?, ?, ?, ?> <?, Warm, ?, ?, ?, ?>}



G_3 : {<Sunny, ?, ?, ?, ?, ?> <?, Warm, ?, ?, ?, ?> <?, ?, ?, ?, ?, Same>}

Training Example:

4.<Sunny, Warm, High, Strong, Cool, Change>, EnjoySport = Yes